


Estimated training FLOPs vs N (k=32, sparse top-32 teacher)



Log-log slope: Dense 1.96, Exact top-k 1.80, ARIAD 1.00.

Estimated arithmetic FLOPs (cost model, not measured); excludes softmax, top-k, dedup and memory traffic.