



Training dynamics (N=10,000, k=32, sparse top-32 teacher)

(a) The neighbor graph changes during training



(b) MSE convergence



(a) Each method's own 32-neighbor graph S_t in its own run, audited on all N rows every epoch (not part of training cost); replacement = $1 - |S_t \cap S_{t-1}| / 32$, averaged over rows. Dense attends to all nodes; its line tracks the top-32 of its learned scores. ARIAD's early peak is its first refreshes away from the random initial graph.
(b) MSE at epoch 1 and every 10 epochs. Both panels come from the same seed-0 training runs (full batch, 1 step = 1 epoch, Adam lr=1e-2, identical initialisation for all methods).