


MSE convergence (N=10,000, k=32, sparse top-32 teacher)



MSE recorded every 10 epochs; full-batch training (1 step = 1 epoch).
Adam lr=1e-2; same seed and initialisation for all methods.